

TraceLens: Multimodal Relative Debugging of UI Test Regressions via Pass–Fail Trace Comparison

Musab Ahmed Khan*

Sabanci University

Istanbul, Turkey

musab.khan@sabanciuniv.edu

ABSTRACT

When an automated UI test regresses from pass to fail, engineers practice *relative debugging*: they compare two executions of the same test to find the earliest step whose divergence explains the failure, separating root cause from downstream symptoms. We present TRACELENS, a multimodal pipeline that aligns pass and fail traces, scores per-step *relative anomalies* across network logs, console output, actions, navigation, and screenshots, and optionally reranks candidates with a large language model (LLM) and a vision-language model (VLM). Our first architecture (v1) stacked heuristic ranking, LLM reranking, VLM ensemble, and trace-specific Hit@1 guards; hybrid mode proved brittle when guards tuned for one failure pattern regressed others. We therefore introduce v2, which replaces post-hoc overrides with *unified weighted fusion* and read-only diagnosis. On a benchmark of 22 real UI regression traces, a heuristic baseline achieves 64% Hit@1. LLM-only reranking reaches **86% Hit@1**, the best top-1 localization. Hybrid LLM+VLM fusion reaches **82% Hit@1** but **100% Hit@3** with lower mean rank distance (0.27 vs. 0.32), recovering visually dominant cases while occasionally regressing when visual noise overrides causal text. We recommend LLM-only for sharp top-1 localization and hybrid fusion when top-*k* robustness and screenshot evidence matter.

CCS CONCEPTS

• **Software and its engineering** → **Software testing and debugging**; • **Computing methodologies** → *Machine learning*.

KEYWORDS

relative debugging, regression diagnosis, UI testing, execution traces, fault localization, LLM, VLM

1 INTRODUCTION

Automated end-to-end (E2E) UI tests produce rich *execution traces*: ordered steps recording browser actions, HTTP traffic, console messages, and screenshots. When a test that passed yesterday fails today, debugging is inherently *relative*: engineers diff the good run against the bad run to locate where behavior first diverged. We call this task **relative debugging**, following the compare-against-reference paradigm of Abramson et al. [1, 11]: given a passing trace *P* and a failing trace *F* for the same test case, rank execution steps by how likely each is the *earliest root cause* of the regression.

Relative debugging differs from classical source-level fault localization [2, 13]. The artifact under analysis is an *execution trace*, not a program graph. The supervision signal is the pass–fail *pair* itself: anomalies are defined relative to what changed between runs.

Three difficulties recur in practice. **(1) Symptom vs. cause.** A wrong credential at step 7 may surface only as a verification failure at step 8; ranking by error magnitude favors the symptom [9]. **(2) Silent and visual faults.** Missing form fields, CAPTCHA walls, and stale CDN assets may produce no console or network signal; only screenshots reveal the divergence [8]. **(3) Noise.** Telemetry URLs, placeholder domains, and transient network flicker create spurious diffs that distract both heuristics and language models.

We built TRACELENS to automate relative debugging for UI test traces. The system parses heterogeneous trace formats, aligns pass and fail steps, extracts multimodal relative features, ranks suspicious steps, and generates technical and stakeholder-facing explanations. We evaluate four modes against a **heuristic baseline** (pass–fail signal diffing without API calls): LLM reranking for causal text reasoning, VLM analysis for screenshot pairs, and hybrid fusion combining both.

Running example. On the hackernews trace, a Firebase timeout at step 12 causes many quiet steps afterward; heuristics rank a later verify step with higher pixel score. The LLM reranker reads `errors_observed_on_next_step` and promotes step 12, a canonical *relative* causal pattern. On amazon, the ground-truth fault is an auth-gated click (step 20); LLM correctly places it first, but hybrid fusion elevates an earlier scroll step with large screenshot diff, illustrating the LLM/hybrid tradeoff we quantify in Section 4.

Our evaluation on 22 traces shows complementary strengths rather than a single winner. The heuristic baseline reaches 64% Hit@1. **LLM-only** mode achieves the highest Hit@1 (**86%**), indicating strong text and causal ranking when logs and actions are informative. **Hybrid** fusion matches or exceeds LLM on Hit@3 (100% vs. 95%) and on rank-quality metrics (lower mean rank distance and MAD@5), recovering visual-dominant or ambiguous cases (e.g., github) while occasionally regressing when visual noise overrides text (e.g., amazon). We therefore treat LLM as the default for sharp top-1 localization and hybrid as the multimodal setting when top-*k* robustness and visual evidence matter.

Contributions.

- **TraceLens**, an end-to-end relative debugging pipeline for UI regression traces with interpretable multimodal signals and stakeholder reports.
- **Multimodal relative features**, including navigation mismatch detection, visual-causal walk-back from persistent screenshot divergence, and observer vs. action attribution.
- **v2 unified fusion**, replacing v1’s trace-specific guard stack with a single weighted ranker; LLM/VLM outputs are soft priors and diagnosis is read-only.

*CS 460 Project, Sabanci University. Student ID: 33017.

- **Empirical study** on 22 diverse traces with ablations and an LLM vs. hybrid tradeoff analysis.

2 BACKGROUND AND RELATED WORK

Relative and delta debugging. Abramson et al. [1] introduced *relative debugging*: locating errors by comparing a suspect program’s state against a trusted reference execution. The GUARD tool [11] operationalized this for scientific software evolution. We adapt the same principle to UI test traces: the pass run is the reference and the fail run is the suspect; divergence is measured per aligned step rather than per variable. Delta debugging [15] isolates failure-inducing *inputs* by systematically simplifying test cases; relative debugging instead ranks *execution steps* where pass and fail behavior diverge. Trace comparison also appears in log analysis [7, 9] and web testing [8, 10], but prior tools rarely publish reproducible ranking metrics over multimodal UI traces with ground-truth fault steps.

Fault localization. Spectrum-based and learning-based fault localization [12, 13] rank *source entities* (statements, methods) using coverage and pass/fail test outcomes. TRACELENS instead ranks *trace steps* using direct pass–fail diffs of actions, logs, and pixels. The tasks are complementary: trace-step localization supports test maintainers; code-level FL supports developers patching production code.

LLMs and VLMs for software engineering. Recent surveys document LLM use for bug repair, log parsing, and code generation [3, 5]. We use an LLM as a *reranker* over compact relative features, not as a patch generator. Vision-language models enable GUI understanding [4, 6, 14]; we apply VLMs to *pass vs. fail screenshot pairs* at candidate steps identified by cheaper heuristics, controlling API cost.

Gap. No prior work combines (i) heterogeneous real UI trace schemas, (ii) explicit relative multimodal features, (iii) LLM and VLM priors in one fusion ranker, and (iv) Hit@ k evaluation including text-invisible faults, in a reproducible open-source implementation with frozen benchmark runs.

3 APPROACH

Figure 1 summarizes our v2 architecture (entry point `main2.py`). Section 3.6 briefly describes v1, which we supersede.

3.1 Task formulation

Input: passing trace $P = (p_0, \dots, p_{n-1})$ and failing trace $F = (f_0, \dots, f_{m-1})$ for the same test.

Output: ranked step list $(s_1, \dots, s_k)$, a technical diagnosis, and a plain-language summary.

Each step is unified into a schema with action text, action type, optional intent field (one data source), network logs, console logs, and screenshot path. Alignment uses index-based pairing when $|P| = |F|$; otherwise action-similarity alignment when faults insert or delete steps.

Ground truth is a single fault step index per trace (used only for evaluation). Metrics: Hit@ k (GT in top- k), mean rank distance $|\text{rank}(GT) - 1|$, and MAD@5 (mean absolute deviation of GT rank from 1, capped at 5).




Figure 1: TRACELENS v2 pipeline. Pass and fail traces are parsed, aligned, and differenced into fusion channels; optional LLM/VLM priors feed a single weighted sort. Diagnosis is read-only.

Table 1: Pass–fail relative signals in TRACELENS.

Signal	Relative comparison	Example
Network	New/changed errors in fail	HTTP 500
Console	New JS errors in fail	UI crash
Action	Text or type diverges	Invalid form
Navigation	Wrong page loaded	Wrong click
Pixel	Screenshot diff	Visual regression

3.2 Relative signal extraction (heuristic baseline)

The **heuristic baseline** scores each aligned pair (p_i, f_i) without API calls. Four text signals in $[0, 1]$ are combined with fixed weights (renormalized when intent is absent):

- **Network:** new HTTP errors, status changes, or missing expected requests in F .
- **Console:** new JavaScript errors or warnings in F .
- **Action:** divergence in action text or type between P and F .
- **Intent:** changed verification outcome strings (available on one source schema).

Navigation signals detect wrong-page loads: expected URL absent and unintended page present (`wrong_navigation`), boosting steps that explain silent navigation faults.

Noise filtering removes placeholder domains (e.g., `example.com` distractors) and known telemetry patterns so relative diffs focus on test-relevant traffic.

Pixel diff (local, no API) compares pass/fail PNGs for the heuristic top-5; enabled by default, disabled with `-no-pixel`.

Table 1 summarizes signals; the baseline mode is `python main2.py` without flags.

3.3 Visual relative signals

Three visual layers differ in scope and cost:

Pixel boost computes mean pixel difference on aligned PNG pairs (local, cheap).

Visual-causal attribution scans all steps when visual mode is enabled, finds the earliest *persistent* screenshot divergence, and walks back to a silent cause step (e.g., missing zipcode visible only on the next screen). This uses deterministic rules, not a neural model.

VLM analysis sends pass/fail pairs for top- K candidates (~ 5 steps) to a vision-language model, returning per-step visual scores and a suggested root step. Visual-causal steps are injected into the VLM candidate pool; scores can backfill from symptom steps onto cause steps.

3.4 LLM relative reranking

With `-llm`, a language model reranks a candidate pool: heuristic top-15 plus injected steps with strong navigation, causal, or action-changed signals. The prompt exposes compact *relative* fields (action_changed, wrong_navigation, filtered_network/console snippets) without screenshots.

The LLM returns an ordered list and a diagnosis. In `v2`, rerank positions map to an `llm_prior` fusion channel; diagnosis **does not** change the final rank.

We use OpenRouter with temperature 0.0 (Nemotron-3-Super-120B for text) to maximize reproducibility; provider routing may still introduce variance.

3.5 v2 unified fusion

`v2`'s main design contribution is a **single fusion score** per step:

$$S(i) = \sum_{c \in C} w_c \cdot f_c(i) \quad (1)$$

where channels C include text, navigation, causal action, observer symptom, visual causal, pixel, and optionally LLM and VLM priors (denoted `llm_prior` and `vlm_prior` in code). Weights w_c are fixed in configuration (`llm_prior = 0.28`, `vlm_prior = 0.18`, mirroring `v1`'s 60/40 LLM/VLM blend).

Causal attribution distinguishes *action* steps (clicks, typing) from *observer* steps (verify, wait). When errors appear on a verify step, a preceding action step may be the true relative root (causal_action channel). When the verify step itself is labeled as ground truth (symptom-on-verify faults), observer channels must not be over-capped; `v1` learned this through per-trace guards, `v2` encodes it as channel weights and the observer-after-causal cap.

Structural policies (not trace-specific guards) adjust scores:

- **Observer cap:** limit verify-step visual weight when a prior causal action step has strong text support.
- **Downstream penalty:** demote steps after inferred root, with exemption when text signal exceeds a threshold.
- **LLM/action bonuses:** small nudges when the LLM ranks a causal or action-changed step first.

Steps are sorted once by $S(i)$. Hybrid mode (`-llm -vlm`) activates all channels; LLM-only and VLM-only use the same scorer with different channel subsets.




Figure 2: v1 stacked pipeline (main.py). Highlighted stages can override prior rankings. Hybrid runs LLM overrides before VLM overrides and a second guard pass.

Mode summary.

- **Heuristic (baseline):** fast, interpretable pass-fail diff; 64% Hit@1; misses causal reasoning on hard traces.
- **LLM:** best Hit@1 (86%); strong when logs and actions encode causality.
- **VLM:** addresses visual faults; weaker alone on network-heavy traces; higher latency and cost.
- **Hybrid:** best Hit@3 and rank distance; trades one Hit@1 vs. LLM when visual channel overweighted.

3.6 v1 architecture and limitations

Figure 2 shows `v1` (`main.py`), our first production architecture. It validated that relative text and visual evidence are both necessary, but ranking policy was **fragmented**: heuristic sort, then LLM hard rerank, Hit@1 guard, deterministic promote, diagnosis promote, VLM ensemble, and post-blend guard.

Guards encoded patterns from individual benchmark traces (promote causal action on hackernews; lock strong verify-pixel on youtube; suppress spurious scroll on amazon). Tuning for one pattern regressed others; e.g., blocking deterministic promotes under visual lock sharply dropped accuracy on action-labeled traces in an intermediate experiment.

Hybrid mode (`-llm -vlm`) activated both override chains, so conflicts surfaced more often than in single-modality modes. On partial benchmark runs, `v1` hybrid underperformed LLM-only despite each modality working well in isolation. Diagnosis and deterministic rules could also promote steps to rank 1 using patterns discovered while inspecting labeled traces, complicating claims of a general ranking policy.

v2 response: one fusion function, soft LLM/VLM priors, read-only diagnosis, structural caps replacing per-trace guard exceptions. v1 remains in the repository as a reference implementation; all reported main results use v2.

4 EVALUATION

4.1 Research questions

RQ1: How much do LLM and multimodal fusion improve relative step ranking over the heuristic baseline?

RQ2: What is the contribution of pixel/visual channels?

RQ3: How do LLM-only and hybrid compare on Hit@1 vs. top- k and rank-quality metrics?

RQ4: How does performance vary across data sources and fault signal types?

4.2 Benchmark and setup

We evaluate on **22 traces** from three independent collections: 8 from efe/irem, 10 from areeb/salem, and 4 from ersel (reported as a *test set* in figures due to distinct schema and intent fields). Ground truth specifies one fault step per trace covering network errors, authentication failures, console crashes, screenshot-only faults, absence-of-activity faults, and form/input faults. Table 2 lists a representative subset; the full catalog is in our repository.

All main results use v2 (main2.py) with frozen run JSONs listed in scripts/metrics_manifest.yaml. Baseline: python main2.py. LLM: Nemotron-3-Super-120B via OpenRouter ($T=0$). VLM: Llama-4-Scout-17B via Groq ($T=0$), ~5 screenshot pairs per trace.

4.3 Results

Table 3 reports aggregate v2 metrics on all 22 traces.

RQ1. LLM improves Hit@1 from 64% to 86% (+22 points); hybrid reaches 82% (+18). Every mode except no-pixel heuristic beats the baseline on Hit@3 or rank distance.

RQ2. Removing pixel changes Hit@1 from 64% to 68% on this corpus; marginal on aggregate, but pixel and visual-causal channels are essential for screenshot-primary faults (Section 4.4).

RQ3. Figure 3 summarizes the LLM vs. hybrid tradeoff. Hybrid wins Hit@3 (100% vs. 95%) and mean rank distance (0.27 vs. 0.32) but loses one Hit@1 (18/22 vs. 19/22).

Figure 4 shows Hit@ k across all modes; Figure 5 summarizes multiple metrics on a common scale. Figure 6 gives per-trace Hit@1; Figure 7 breaks Hit@1 by source collection.


4.4 Case studies

amazon (GT step 20, auth-gated action). LLM ranks the causal action first (Hit@1). Hybrid prefers an earlier scroll step with stronger pixel signal (rank distance 2). Illustrates visual noise overriding causal text.

github (GT step 23, JS crash). Both LLM and hybrid miss Hit@1, but hybrid places GT at rank 2 (distance 1) vs. LLM rank 5 (distance 4). Hybrid recovers a visually evident but text-ambiguous fault in the shortlist.

saucedemo_2, **bbc**, **pypi**, **youtube** (screenshot-primary). The heuristic baseline misses saucedemo_2 (no text signal); LLM and

v2 — LLM vs Hybrid: complementary metrics (22 traces)



LLM wins Hit@1 on amazon; Hybrid wins Hit@3 on github and has lower rank distance / MAD@5.

Figure 3: LLM vs. hybrid on Hit@1/3, mean rank distance, and MAD@5 (v2, 22 traces).




Figure 4: Hit@1/3/5 across modes; v1 vs. v2 evolution. Heuristic is the baseline.




Figure 5: Multi-metric radar comparison across v1 and v2 modes (Hit@1/3/5, rank distance, MAD@5).

hybrid both localize it using local visual-causal and pixel channels in v2.



Table 2: Representative benchmark traces (22 total). GT = ground-truth fault step (0-indexed).

Trace	Source	GT	Fault type	Primary signal
hackernews	areeb/salem	12	Firebase timeout	Console + network
amazon	areeb/salem	20	Auth-gated action	Console + action
github	areeb/salem	23	JS rendering crash	Console
saucedemo_2	efe/irem	27	Missing zipcode	Screenshot only
wolfram	efe/irem	12	Wrong element clicked	Absence of activity
wikipedia	efe/irem	13	Broken anchor / TOC	Network noise
opencart_wishlist_35	ersel	22	HTTP 404 product	Intent + network

Table 3: v2 aggregate results on 22 traces. Baseline = heuristic with pixel. VLM $n=21$ (dictionary excluded: all-zero visual scores).

Mode	n	Hit@1	Hit@3	Hit@5	MRD
Heuristic	22	64%	91%	91%	0.86
No pixel	22	68%	86%	91%	0.86
LLM	22	86%	95%	100%	0.32
VLM	21	67%	95%	100%	0.43
Hybrid	22	82%	100%	100%	0.27

MRD = mean rank distance. MAD@5: Heuristic 7.94, LLM 6.94, Hybrid 6.07.

**Figure 7: v2 Hit@1 by source collection and mode.****Figure 6: Per-trace Hit@1 heatmap (v2). Left band: source collection (ersel = test set).**

4.5 Analysis by signal type

Grouping traces by primary signal type reveals expected mode strengths. **Network/console faults** (9 traces): LLM reaches 89% Hit@1 (8/9) vs. 67% (6/9) for the heuristic baseline. **Screenshot-primary faults** (4 traces): the baseline reaches 75% Hit@1 (3/4); both LLM and hybrid reach 100% because v2’s local visual-causal scan runs even in LLM-only mode. **Absence-of-activity** (wolfram, both wikipedia traces): navigation and missing-page features are critical; LLM reranking promotes the silent navigation fault over CDN noise.

4.6 v1 comparison

On partial runs, v1 heuristic reached 55% Hit@1 ($n=22$); v1 hybrid $\sim 70\%$ ($n=10$) did not reliably beat v1 LLM ($\sim 72\%$, $n=18$) due to guard interactions. v2 hybrid on the full corpus reaches 82% Hit@1 with 100% Hit@3 (Table 4 in the appendix).

5 THREATS TO VALIDITY

Construct validity. Single fault-step labels simplify evaluation but blur symptom vs. cause (e.g., amazon). Hit@ k and rank distance partially mitigate this by rewarding near-misses.

Internal validity. Fusion weights were fixed in v2/config.yaml before aggregate reporting. Structural rules were motivated by recurring failure *types*, not by per-trace grid search on ground truth in v2.

External validity. 22 English UI traces from academic collectors are not industry-scale. Three trace schemas limit generalization without adapter code.

Conclusion validity. LLM/VLM APIs with $T=0$ may still vary by provider routing; we report frozen run timestamps in the repository.

Evaluation bias. ersel (4 traces) uses intent fields absent elsewhere; VLM aggregate excludes dictionary ($n=21$).

6 CONCLUSION

Relative debugging of UI test regressions benefits from multimodal pass-fail comparison. TRACELENS demonstrates a reproducible pipeline from aligned trace diffs through LLM/VLM priors to ranked suspects and stakeholder explanations. Against a 64% heuristic baseline, LLM-only achieves 86% Hit@1; hybrid fusion trades one top-1 hit for perfect Hit@3 and better rank quality. v2’s unified fusion replaces v1’s trace-specific guard brittleness with a single interpretable scorer.

Future work includes a larger public relative-debugging benchmark, learned fusion weights, multi-fault labels, CI integration, and practitioner user studies.

DATA AVAILABILITY

The repository includes trace data, ground truth, the v2 pipeline (main2.py), frozen metric JSONs, and figure scripts. Regenerate figures with `python scripts/plot_metrics.py`.

A V1 AGGREGATE COMPARISON

Table 4: v1 vs. v2 Hit@1 on overlapping runs.

Mode	v1 Hit@1	v2 Hit@1
Heuristic	55% ($n=22$)	64% ($n=22$)
LLM	72% ($n=18$)	86% ($n=22$)
VLM	88% ($n=17$)	67% ($n=21$)
Hybrid	70% ($n=10$)	82% ($n=22$)

v1 VLM headline accuracy on 17 traces should not be compared directly to v2 VLM on 21 traces; different coverage and fusion integration dominate.

REFERENCES

- [1] David Abramson, Ian Foster, John Michalakes, and Rok Sosić. 1996. Relative Debugging: A New Paradigm for Debugging Scientific Applications. *Commun. ACM* 39, 11 (1996), 67–77.
- [2] Hiralal Agrawal, Joseph R. Horgan, and Si-Leung Wong. 1995. Fault Localization Using Execution Slices and Dataflow Tests. *IEEE Transactions on Software Engineering* 21, 2 (1995), 149–162.
- [3] Placeholder Authors. 2024. Large Language Models for Software Engineering: A Systematic Literature Review. In *Proceedings of the Conference Name Placeholder (Venue 'XX)*. Replace with published survey before submission.
- [4] Placeholder Authors. 2024. Vision-Language Models for Graphical User Interface Understanding. In *Proceedings of the Conference Name Placeholder (Venue 'XX)*. Replace with WebArena/Mind2Web or related citation.
- [5] Tom B. Brown et al. 2020. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [6] Xiang Deng et al. 2023. Mind2Web: Towards a Generalist Agent for the Web. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [7] Mark Grechanik, Chen Fu, and Qiang Xie. 2009. Automatically Finding Performance Problems with Feedback-Directed Monitoring as You Go. In *Proceedings of the International Symposium on Software Testing and Analysis (ISSTA)*. 227–238.
- [8] Ali Mesbah et al. 2012. Crawling Ajax-Based Web Applications through Dynamic Analysis of User Interface State Changes. In *ACM Transactions on the Web*, Vol. 6. 1–30.
- [9] Alessandro Panichella et al. 2015. How Developers Debug Software: The Role of Log Points. In *Proceedings of the International Conference on Software Maintenance and Evolution (ICSME)*. 1–10. Log-based debugging representative.
- [10] Filippo Ricca and Paolo Tonella. 2001. Analysis and Testing of Web Applications. In *Proceedings of the International Conference on Software Engineering (ICSE)*. 25–34.
- [11] Rok Sosić and David Abramson. 1997. GUARD: A Relative Debugger. *Software: Practice and Experience* 27, 2 (1997), 185–206.
- [12] Ming Wen et al. 2020. Historical Spectrum Based Fault Localization. In *Proceedings of the International Conference on Software Engineering (ICSE)*. 1–12. Spectrum-based FL representative.
- [13] W. Eric Wong et al. 2016. A Survey on Software Fault Localization. *IEEE Transactions on Reliability* 65, 3 (2016), 841–867.
- [14] Shunyu Yao et al. 2024. WebArena: A Realistic Web Environment for Building Autonomous Agents. In *International Conference on Learning Representations (ICLR)*.
- [15] Andreas Zeller. 2002. Simplifying and Isolating Failure-Inducing Input. In *IEEE Transactions on Software Engineering*, Vol. 28. 183–200.